

# Understanding common concurrency patterns with tricky examples

---


CppIndiaCon 2023, August 4th 2023

Venkata Naga Ravikiran Bulusu

# About me

---

- Embedded SW developer at VITES GmbH, Munich, Germany
- SW design/development for satellite communication systems
- C/C++17/Embedded Linux
- Contact:
  - Gmail: [ravi.has.kiran@gmail.com](mailto:ravi.has.kiran@gmail.com)
  - LinkedIn: <https://www.linkedin.com/in/venkata-naga-ravikiran-b-0b479b25/>



---


# Agenda

---

- Basics of Concurrency
- Concurrency patterns
- Classic problems and solutions
  - Sleeping barber
  - Dining philosophers
  - Reader-writer
- Asynchronous I/O
- Reactor pattern (libuv)
- Final thoughts
- References

# Basics of concurrency I

- Concurrency is a technique that allows multiple processes to run at the same time by managing access to shared resources on a single CPU core. This is done through interleaved running of processes via context switching.
- One way to achieve concurrency is through **threads**, which are lightweight processes that can switch quickly and share information easily between them.



# Basics of concurrency II

---

## Advantages

- Improved throughput and efficiency
- Better CPU utilization
- Concurrent access for multiple users
- Real-time applications

## Challenges

- Race conditions
- Deadlocks
- Starvation
- Livelock

# Concurrency Patterns

---

## Concurrent Architecture

- Active object
- Monitor object
- Reactor

## Synchronization Patterns

- Dealing with sharing
- Copied value
- Thread-specific storage
- Future

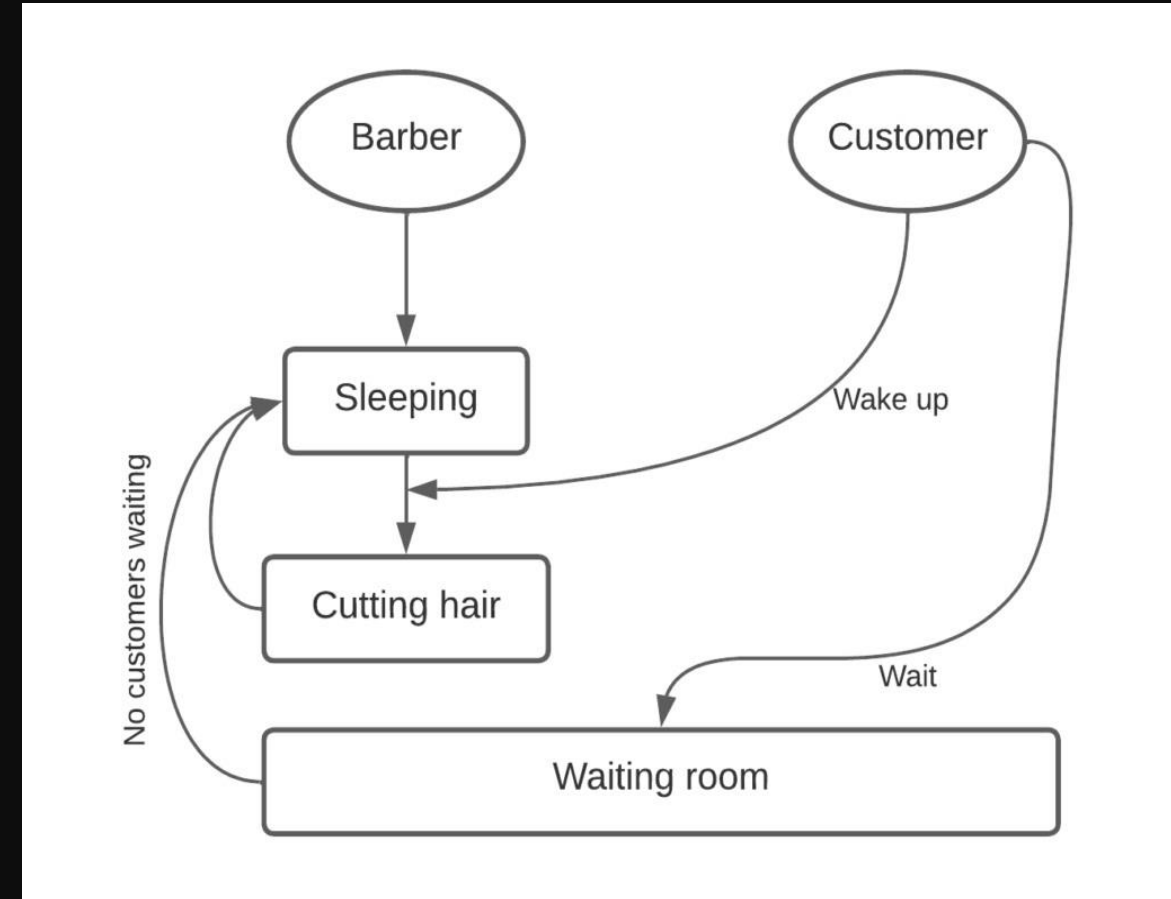
## Dealing with Mutation

- Scoped locking
- Strategized locking
- Thread-safe Interface
- Guarded suspension

# Sleeping Barber

[https://en.wikipedia.org/wiki/Sleeping\\_barber\\_problem](https://en.wikipedia.org/wiki/Sleeping_barber_problem)

- Monitor object: A pattern that encapsulates shared data and its synchronization mechanisms inside a class
- Objectives:
  - No race condition
  - No starvation
  - No deadlock
  - Mutual exclusion
  - Efficient resource allocation




# Dining Philosophers

---

[https://en.wikipedia.org/wiki/Dining\\_philosophers\\_problem](https://en.wikipedia.org/wiki/Dining_philosophers_problem)

Objectives:

- No race condition
  - No starvation
  - No deadlock
  - Mutual exclusion
  - Efficient resource allocation
- 





# Readers-writers


[https://en.wikipedia.org/wiki/Readers%E2%80%93writers\\_problem](https://en.wikipedia.org/wiki/Readers%E2%80%93writers_problem)

## Objectives:


- Mutual exclusion for writers
- No writer starvation
- No reader-writer starvation
- Multiple readers allowed
- Read-Copy-Update (RCU) consistency (handled only in kernel for most cases)

## Readers-Writers Problem in Operating System

When Readers are accessing the Database



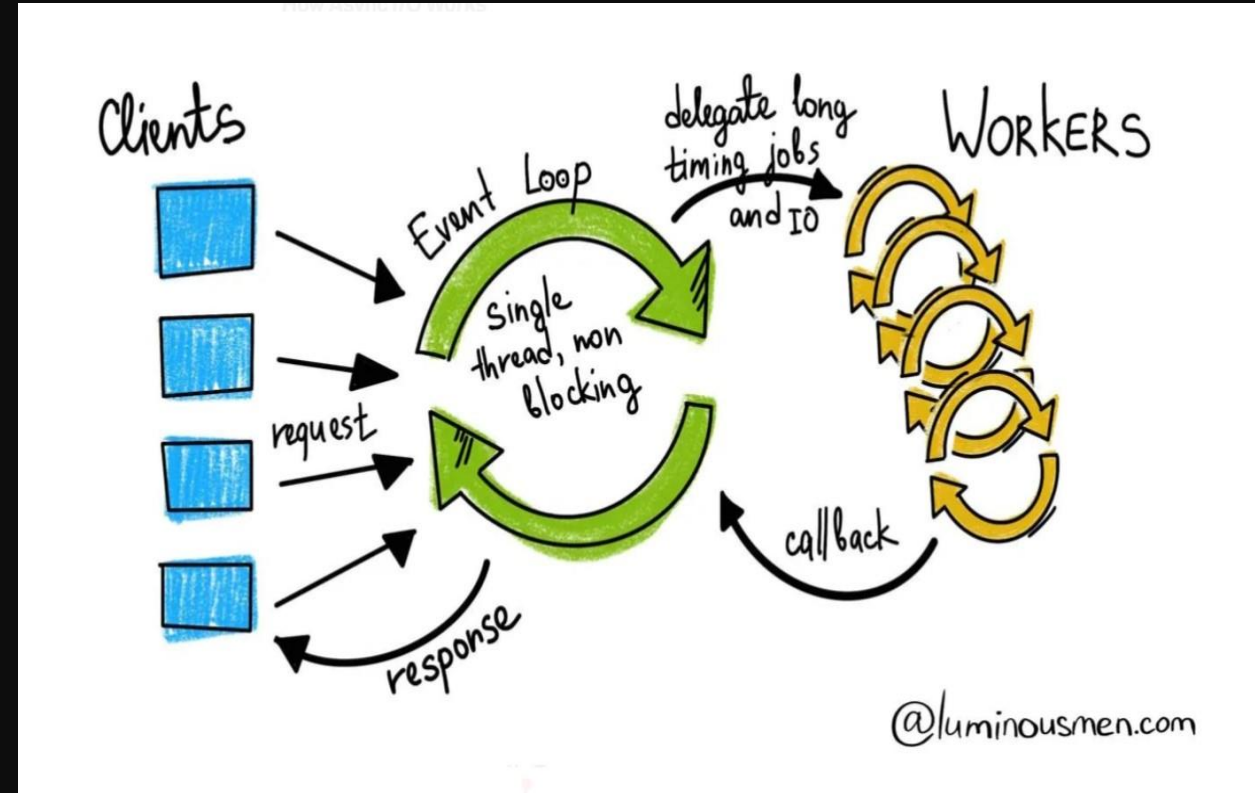
When Writers are accessing the Database



# Asynchronous I/O

**Asynchronous I/O** (also **non-sequential I/O**) is a form of input/output processing that permits other processing to continue before the transmission has finished

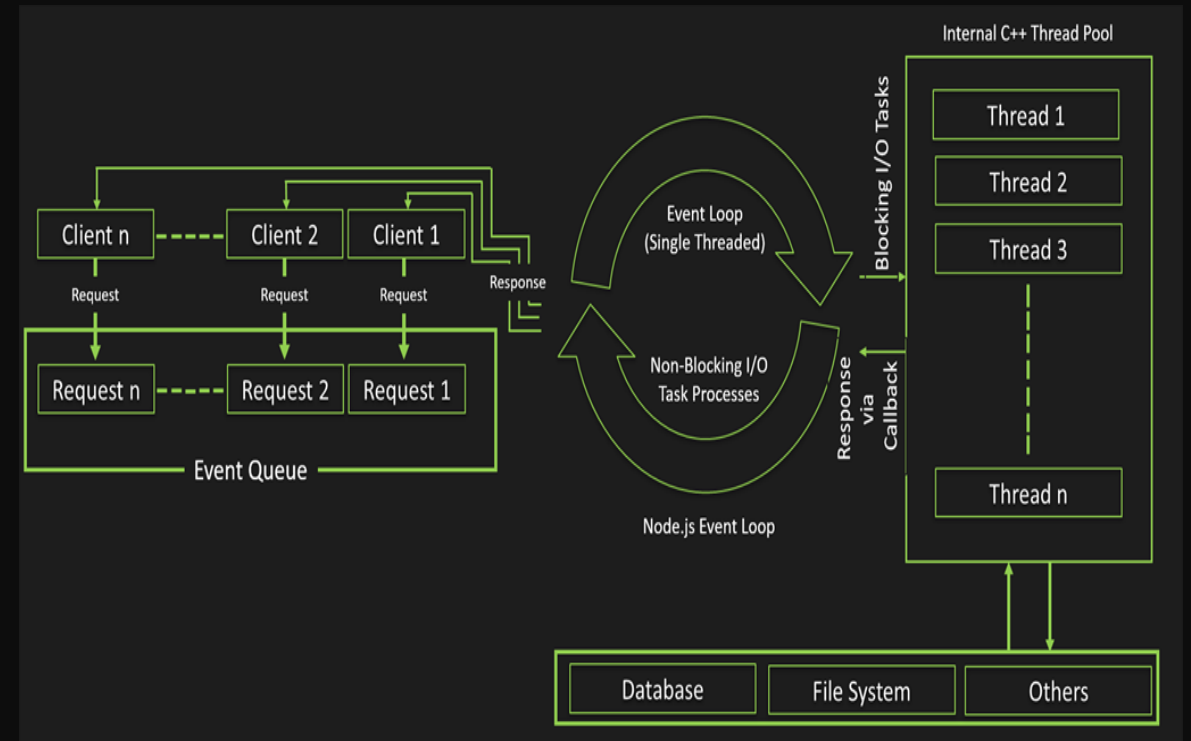
- This provides opportunities for a program to continue running other code while waiting for a long-running task to complete
- The time-consuming task is executed in the background while the rest of the code continues to execute
- <https://libuv.org/>



# Reactor pattern I (libuv)

Key components:

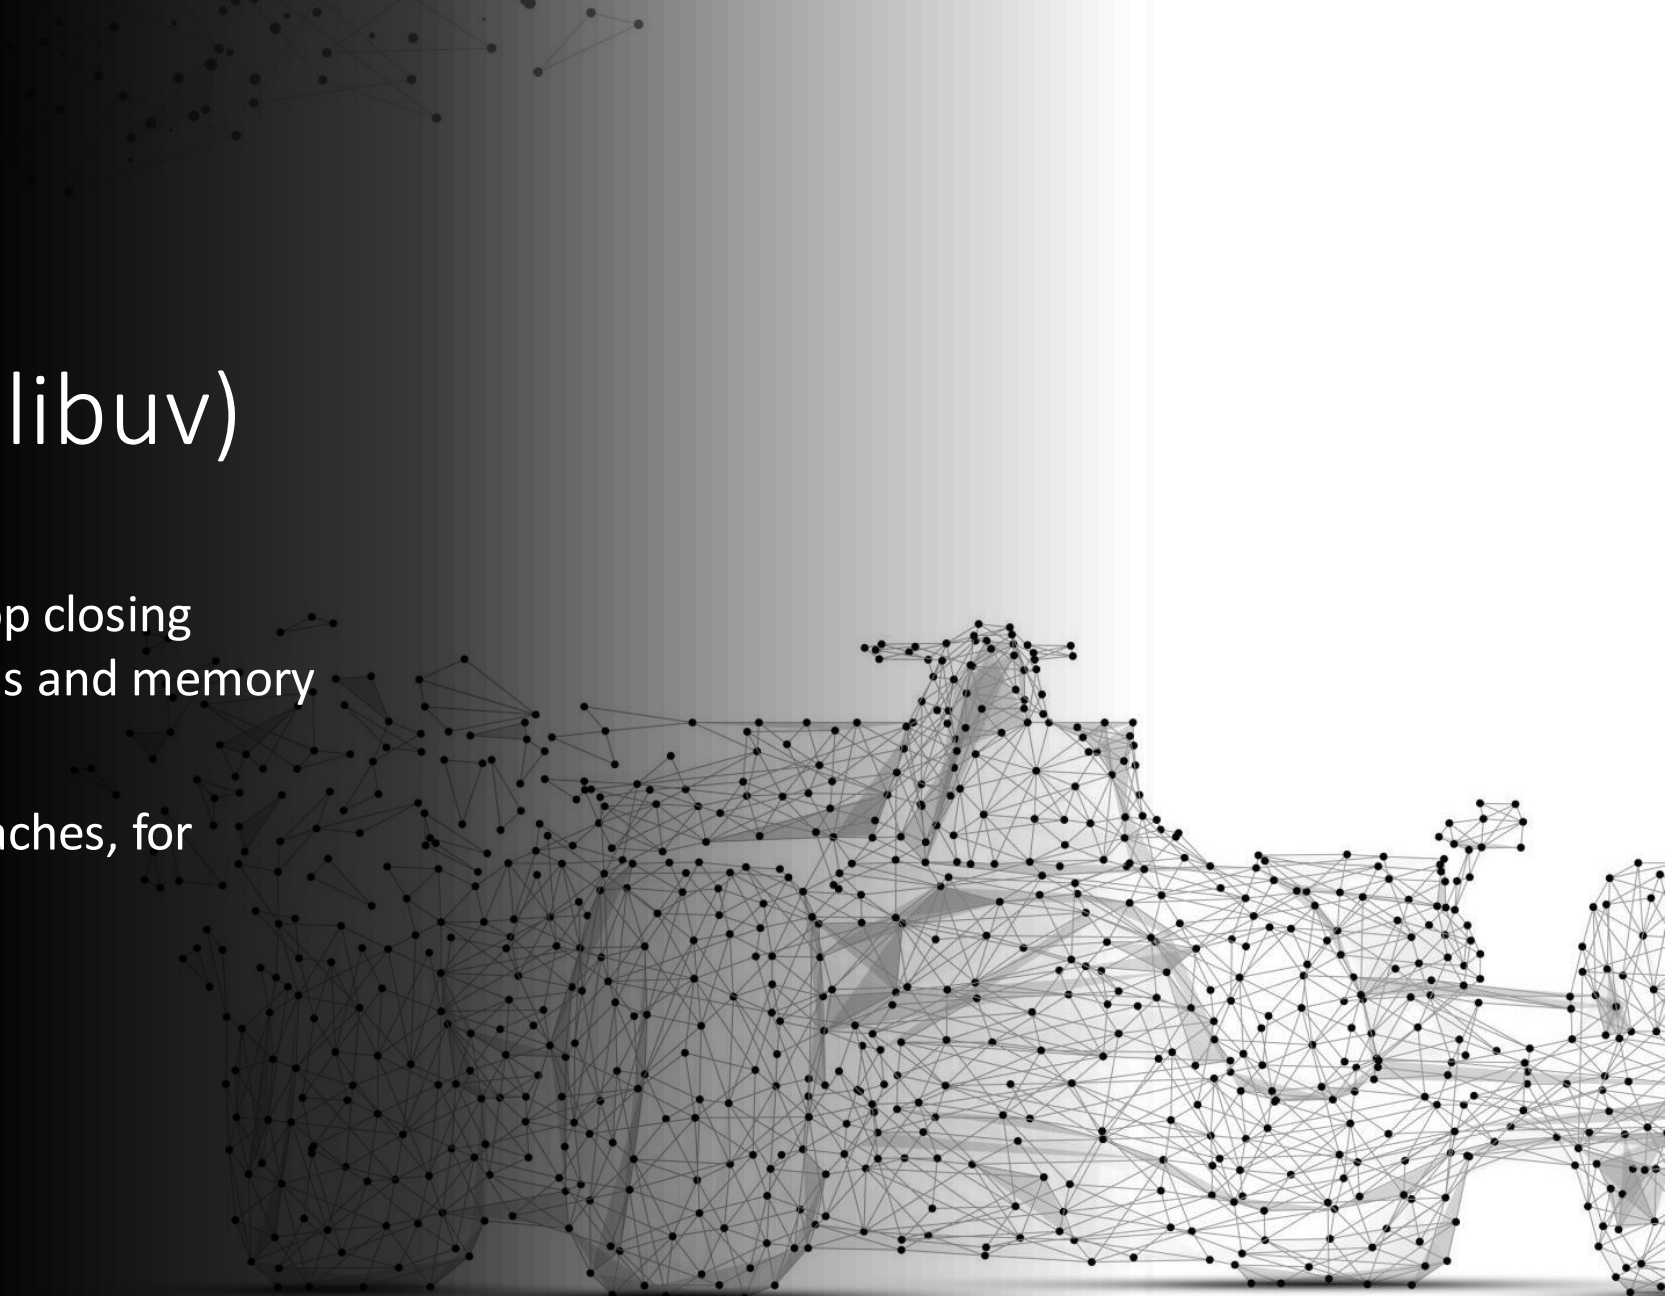
- Event loop -> single threaded
- Event queue -> queue requests
- Execution handlers -> callback response





# Reactor pattern II (libuv)

---

- Problem: Design an libuv event loop closing mechanism without race conditions and memory errors
  - Examples to study different approaches, for solving the problem
- 
- An abstract wireframe sculpture of a horse, composed of numerous black dots connected by thin lines, creating a mesh-like structure. The sculpture is positioned on the right side of the slide, extending from the middle to the bottom. The background is a gradient from dark grey on the left to light grey on the right, with a vertical white line separating the two halves. The overall aesthetic is modern and technical.

# Final Thoughts I

---

## Best practices

- Understand your model
- Minimize shared mutable state
- Design for failure
- Test with realistic workloads
- Identifying critical sections

## Techniques

- Futures and promises
- Reactive programming
- Non-blocking algorithms
- Transactional memory
- Immutable data structures

# Final Thoughts II

---

## Performance considerations

- Measuring and profiling concurrent applications
- Load balancing and work distribution
- Granularity of locks
- Scaling on multi-core processors

At times, concurrency issues are hard to understand and debug. But solving them gives a big chance to step up the ladder in software design

---

# References

---

- [cplusplusreference.com](http://cplusplusreference.com)
  - [https://github.com/CppCon/CppCon2021/blob/main/Presentations/ConcurrencyPatterns\\_1.pdf](https://github.com/CppCon/CppCon2021/blob/main/Presentations/ConcurrencyPatterns_1.pdf)
  - <http://docs.libuv.org/en/v1.x/>
  - [https://en.wikipedia.org/wiki/Reactor\\_pattern](https://en.wikipedia.org/wiki/Reactor_pattern)
  - <https://hashnode.com/post/node-js-event-loop-behind-the-scenes-reactor-pattern-cksds9t6r0eilu5s19gn49oct>
  - <https://valgrind.org/docs/manual/manual.html>
  - <https://levelup.gitconnected.com/multithreading-and-concurrency-concepts-i-wish-i-knew-before-the-interview-11895226179>
  - <https://medium.com/nerd-for-tech/what-is-asynchronous-i-o-b37994359471>
  - <https://medium.com/the-elegant-code/set-another-place-for-plato-55eda132e43b>
  - <https://www.c-sharpcorner.com/article/node-js-event-loop/>
-





THANK YOU!